



UNIVERSITY OF COPENHAGEN

Graph Fundamentals and Graph Neural Networks

NORA Summer School on Geometric Deep Learning

Raghavendra Selvan

Assistant Professor (TT)

Dept. of Computer Science (ML Section)

University of Copenhagen

Pioneer Centre for Artificial Intelligence, Denmark

raghav@di.ku.dk

X /raghavian

<https://raghavian.github.io>

Overview of Lectures

■ Lecture-1

- Motivation
 - Why Graph ML?
 - Deep Learning Basics
- Graph Learning in Spectral domain
 - ChebyNet
- Graph Learning in Spatial domain
 - Node GNNs
- Applications: Node/Graph Classification

■ Lecture-3

- Graph Representation Learning
- Self-Supervised GNNs
- Generative Modelling for Graphs
 - Variational Autoencoders
 - Diffusion Models
- Resource Efficiency and Graph ML
 - Energy Consumption Aware NAS
 - Extreme Quantization of Activation Maps
- Summary

■ Lecture-2

- Learning on Graphs
- GraphSAGE
- Jumping Knowledge Networks
- Graph Isomorphism Network
- Graph Attention Network
- Relational GNNs



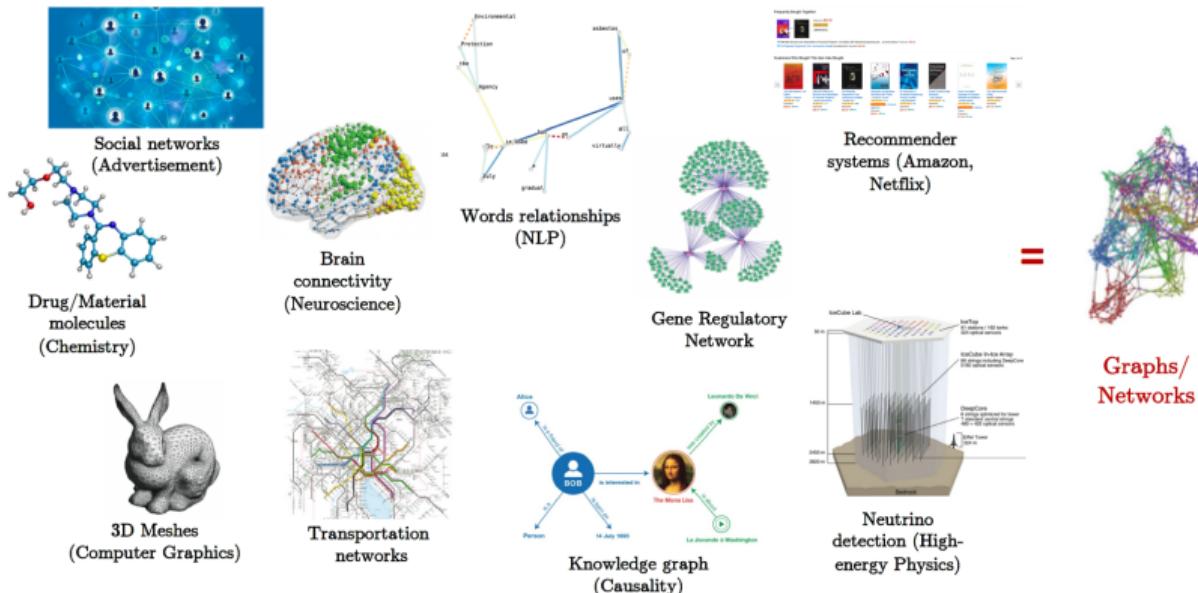
Graph Fundamentals and GNNs: Overview

1 Lecture-1

- Motivation
 - Why Graph ML?
 - Deep Learning Basics
- Graph Learning in Spectral domain
 - ChebyNet
- Graph Learning in Spatial domain
 - Node GNNs
- Applications: Node/Graph Classification



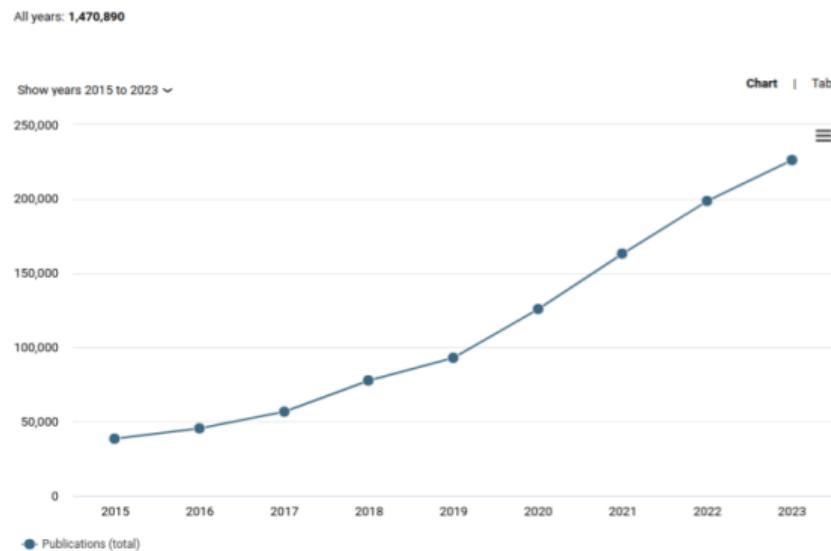
Graphs are everywhere!



Heterogenous graph-structured data are non-Euclidean data



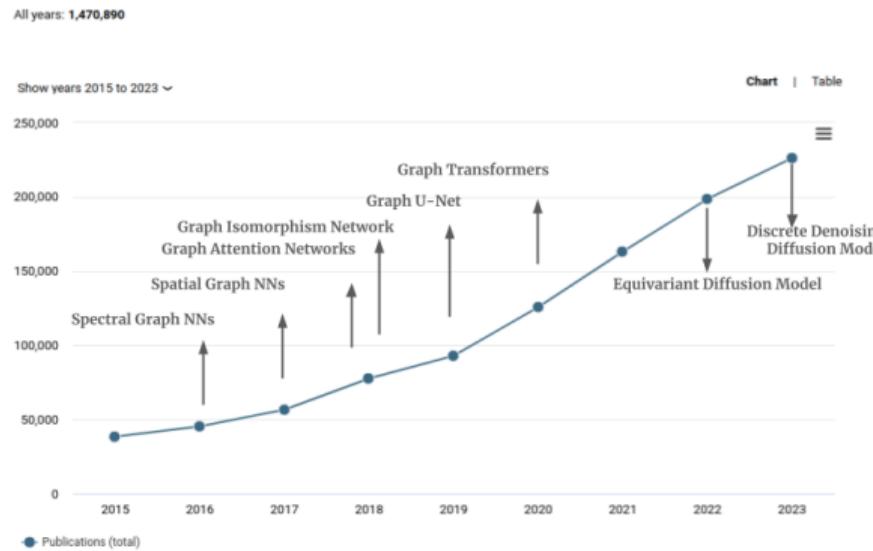
Graph/Geometric Deep Learning is also going everywhere!



-
- Adapted from <https://app.dimensions.ai/>
- Defferrard et al. 2016. Convolutional neural networks on graphs with fast localized spectral filtering.
Kipf & Welling. 2017. Semi-supervised classification with graph convolutional networks.
Veličković et al. 2018 "Graph attention networks."; Xu, Keyulu, et al. 2018 "How Powerful are Graph Neural Networks?
Gao & Ji. 2019. Graph U-Nets; Yun et al. 2019. Graph Transformer Networks
Hoogeboom et al. Equivariant Diffusion for Molecule Generation in 3D. 2022; Vignac et al. DiGress: Discrete Denoising diffusion for graph generation. 2023



Graph/Geometric Deep Learning is also going everywhere!



-
- Adapted from <https://app.dimensions.ai/>
- Defferrard et al. 2016. Convolutional neural networks on graphs with fast localized spectral filtering.
- Kipf & Welling. 2017. Semi-supervised classification with graph convolutional networks.
- Veličković et al. 2018 "Graph attention networks."; Xu, Keyulu, et al. 2018 "How Powerful are Graph Neural Networks?"
- Gao & Ji. 2019. Graph U-Nets; Yun et al. 2019. Graph Transformer Networks
- Hoogeboom et al. Equivariant Diffusion for Molecule Generation in 3D. 2022; Vignac et al. DiGress: Discrete Denoising diffusion for graph generation. 2023



Let us start with a quiz

- Multiple choice questions
- Choose the most correct answer as you see it
- Some specialised terms/acronyms are intentionally used
- Log on to socrative.com
- Choose “Student Login”
- Room name: RAGHAV



One slide introduction to Deep Learning

$$\begin{array}{ll} \text{Observed data} & \mathbf{G} = \{\mathbf{G}_1, \mathbf{G}_2, \dots, \mathbf{G}_N\} \\ \text{Labels / Targets} & \mathbf{Y} = \{\mathbf{y}_1, \mathbf{y}_2, \dots, \mathbf{y}_N\} \\ \text{Decision functions/ Models} & g_{\theta}(\cdot) : \mathbf{G} \rightarrow \mathbf{Y} \end{array}$$

- $g_{\theta}(\cdot)$ can be any DL model
- θ are the trainable parameters; $|\theta| \gg N$
- Over-parameterised, non-linear function approximator
- Supervised training uses large labelled datasets
- Gradient descent based optimisation
- Multi-layered perceptrons, Convolutional/Graph neural networks, Transformers...



Convolutional neural networks (CNNs)

- Learnable convolution kernels
- Technically, not convolution but cross-correlation
- Kernel flipping can be overcome during learning
- Known to learn general image descriptors
- And also, specialized task-specific kernels

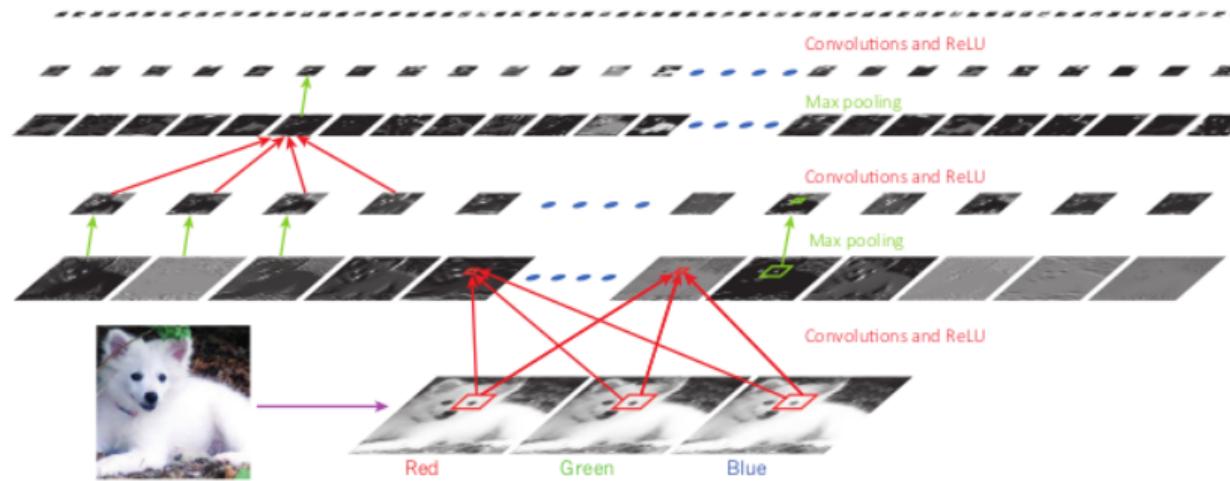


Figure reproduced from Le Cun et al. Deep Learning. 2015



Key properties exploited by CNNs when operating on images, or generally data in Euclidean domain

Success of CNNs in computer vision is due to several properties.

Consider a compact d -dimensional Euclidean domain $\Omega = [0, 1]^d \in \mathbb{R}^d$, on which square integrable functions $f \in L^2(\Omega)$ are defined. Further, consider a translation operation \mathcal{T}

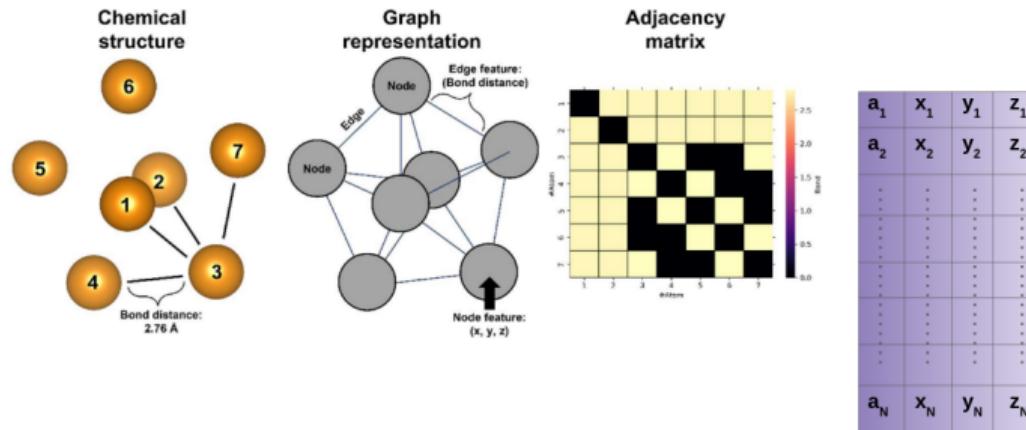
- **Translation invariance:** $y = f(\mathcal{T}(x)) = f(x)$,
- **Translation equivariance:** $y = f(\mathcal{T}(x)) = \mathcal{T}(f(x))$
- **Scale separation:** Long range dependencies from multi-scale interaction terms
- **Compositionality:** $y = f_1 \circ f_2 \circ \dots \circ f_L(x)$
where $f_i(\cdot)$ are comprised of convolution kernels, non-linearities and sometimes pooling operations.

What about permutation invariance: $f(\Pi(x)) = f(x)$ with permutation operator, $\Pi(\cdot)$

$$\int_{-\infty}^{+\infty} |f(x)|^2 dx < \infty$$



Graphs as data structures



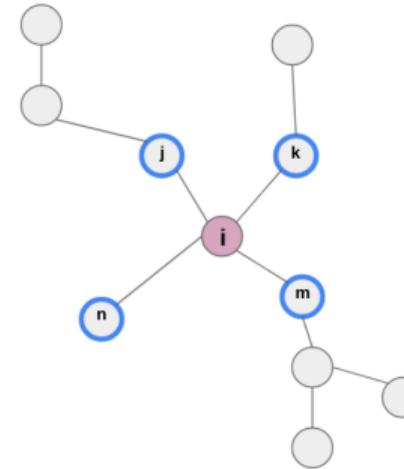
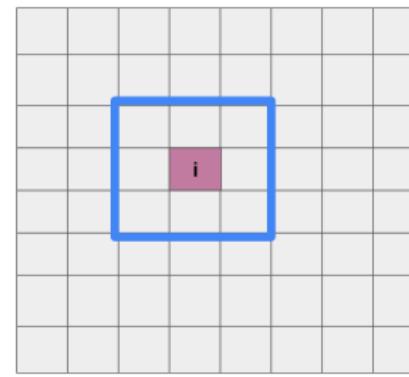
Graph Notations

Consider a graph $\mathcal{G} = (\mathbf{X}, \mathbf{A}, \mathbf{E})$ with N nodes comprising F features per node: $\mathbf{X} \in \mathbb{R}^{N \times F}$ and adjacency matrix $\mathbf{A} \in \{0, 1\}^{N \times N}$ and edge attributes $\mathbf{E} \in \mathbb{R}^{E \times N \times N}$.



Neighbourhoods in Euclidean and non-Euclidean domains

Regular vs Irregular neighbourhoods



Graph Isomorphism



Graph Isomorphism

Given two graphs $G_1 = (V_1, E_1)$ and $G_2 = (V_2, E_2)$, where V_1 and V_2 are the sets of vertices, E_1 and E_2 sets of edges, and A_1 and A_2 are adjacency matrices.



Graph Isomorphism

Given two graphs $G_1 = (V_1, E_1)$ and $G_2 = (V_2, E_2)$, where V_1 and V_2 are the sets of vertices, E_1 and E_2 sets of edges, and A_1 and A_2 are adjacency matrices.

- Two graphs G_1 and G_2 are **isomorphic** if there is a bijection between their vertex sets that preserves adjacency.



Graph Isomorphism

Given two graphs $G_1 = (V_1, E_1)$ and $G_2 = (V_2, E_2)$, where V_1 and V_2 are the sets of vertices, E_1 and E_2 sets of edges, and A_1 and A_2 are adjacency matrices.

- Two graphs G_1 and G_2 are **isomorphic** if there is a bijection between their vertex sets that preserves adjacency.
- Essentially, G_1 and G_2 are structurally identical, even if their vertex labels are different.



Graph Isomorphism

Given two graphs $G_1 = (V_1, E_1)$ and $G_2 = (V_2, E_2)$, where V_1 and V_2 are the sets of vertices, E_1 and E_2 sets of edges, and A_1 and A_2 are adjacency matrices.

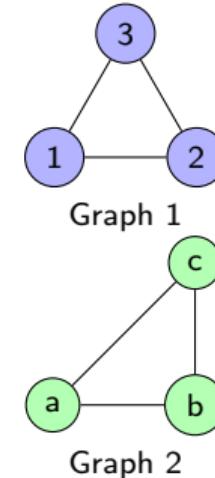
- Two graphs G_1 and G_2 are **isomorphic** if there is a bijection between their vertex sets that preserves adjacency.
- Essentially, G_1 and G_2 are structurally identical, even if their vertex labels are different.
- Problem is to determine whether there exists a bijection $\phi : V_1 \rightarrow V_2$ such that $(u, v) \in E_1$ if and only if $(\phi(u), \phi(v)) \in E_2$ for all $u, v \in V_1$.



Graph Isomorphism

Given two graphs $G_1 = (V_1, E_1)$ and $G_2 = (V_2, E_2)$, where V_1 and V_2 are the sets of vertices, E_1 and E_2 sets of edges, and A_1 and A_2 are adjacency matrices.

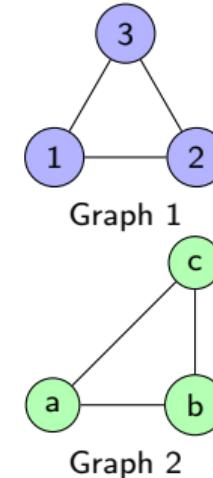
- Two graphs G_1 and G_2 are **isomorphic** if there is a bijection between their vertex sets that preserves adjacency.
- Essentially, G_1 and G_2 are structurally identical, even if their vertex labels are different.
- Problem is to determine whether there exists a bijection $\phi : V_1 \rightarrow V_2$ such that $(u, v) \in E_1$ if and only if $(\phi(u), \phi(v)) \in E_2$ for all $u, v \in V_1$.



Graph Isomorphism

Given two graphs $G_1 = (V_1, E_1)$ and $G_2 = (V_2, E_2)$, where V_1 and V_2 are the sets of vertices, E_1 and E_2 sets of edges, and A_1 and A_2 are adjacency matrices.

- Two graphs G_1 and G_2 are **isomorphic** if there is a bijection between their vertex sets that preserves adjacency.
- Essentially, G_1 and G_2 are structurally identical, even if their vertex labels are different.
- Problem is to determine whether there exists a bijection $\phi : V_1 \rightarrow V_2$ such that $(u, v) \in E_1$ if and only if $(\phi(u), \phi(v)) \in E_2$ for all $u, v \in V_1$.



Formal Definition of Graph Isomorphism

Given two graphs G_1 and G_2 , determine whether there exists a permutation Π of V_1 such that $A_1^\Pi = A_2$, where A_1^Π denotes the adjacency matrix of G_1 under permutation Π .



Many Open Problems in Graph Isomorphism

Complexity



Many Open Problems in Graph Isomorphism

Complexity

- Complexity of the graph isomorphism problem is an area of active research.
- It is not known to be in P or NP-complete, it is widely believed to be in the complexity class NP-intermediate ($P \neq NP$)
- The best known algorithms for graph isomorphism have a complexity of $O(2^{O(\sqrt{n} \log n)})$, where $n = \max(n_1, n_2)$.

Applications



Many Open Problems in Graph Isomorphism

Complexity

- Complexity of the graph isomorphism problem is an area of active research.
- It is not known to be in P or NP-complete, it is widely believed to be in the complexity class NP-intermediate ($P \neq NP$)
- The best known algorithms for graph isomorphism have a complexity of $O(2^{O(\sqrt{n} \log n)})$, where $n = \max(n_1, n_2)$.

Applications

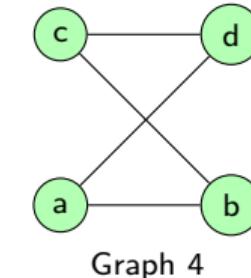
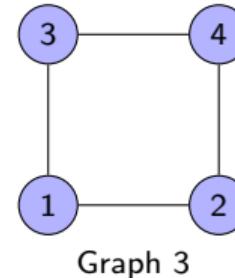
- Chemistry: Identifying molecular structures and chemical compounds.
- Computer Vision: Matching objects and shapes in images.
- Cryptography: Generating cryptographic keys using isomorphic graphs.



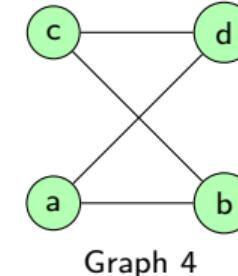
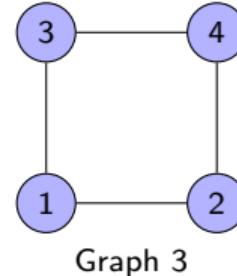
Graph Isomorphism: Example with 4 Vertices



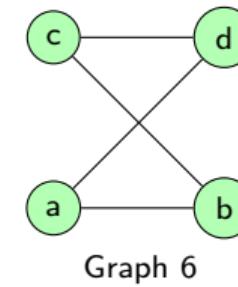
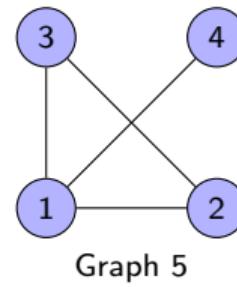
Graph Isomorphism: Example with 4 Vertices



Graph Isomorphism: Example with 4 Vertices



Graph 3 and Graph 4 are isomorphic.



Graph 5 and Graph 6 are not isomorphic.



Complexity of Graph Isomorphism Tests



Complexity of Graph Isomorphism Tests

- **Naive Approach**

- Tries all possible permutations of vertices.
- Complexity: $O(n!)$, where n is the number of vertices.
- Impractical for large graphs due to factorial growth.

- **State-of-the-Art Algorithms**



Complexity of Graph Isomorphism Tests

- **Naive Approach**

- Tries all possible permutations of vertices.
- Complexity: $O(n!)$, where n is the number of vertices.
- Impractical for large graphs due to factorial growth.

- **State-of-the-Art Algorithms**

- Best known exact algorithms: $O(2^{O(\sqrt{n} \log n)})$.
- Exponential but significantly better than the naive approach.

- **Heuristic Algorithms**



Complexity of Graph Isomorphism Tests

- **Naive Approach**

- Tries all possible permutations of vertices.
- Complexity: $O(n!)$, where n is the number of vertices.
- Impractical for large graphs due to factorial growth.

- **State-of-the-Art Algorithms**

- Best known exact algorithms: $O(2^{O(\sqrt{n} \log n)})$.
- Exponential but significantly better than the naive approach.

- **Heuristic Algorithms**

- Various heuristic methods can speed up the process.
- Exploit graph properties and use efficient data structures.

- **Average and Worst-Case Complexity**



Complexity of Graph Isomorphism Tests

- **Naive Approach**

- Tries all possible permutations of vertices.
- Complexity: $O(n!)$, where n is the number of vertices.
- Impractical for large graphs due to factorial growth.

- **State-of-the-Art Algorithms**

- Best known exact algorithms: $O(2^{O(\sqrt{n} \log n)})$.
- Exponential but significantly better than the naive approach.

- **Heuristic Algorithms**

- Various heuristic methods can speed up the process.
- Exploit graph properties and use efficient data structures.

- **Average and Worst-Case Complexity**

- Average-case complexity varies with graph distribution.
- Worst-case complexity is the most critical measure.

- **Implementation Factors**



Complexity of Graph Isomorphism Tests

- **Naive Approach**

- Tries all possible permutations of vertices.
- Complexity: $O(n!)$, where n is the number of vertices.
- Impractical for large graphs due to factorial growth.

- **State-of-the-Art Algorithms**

- Best known exact algorithms: $O(2^{O(\sqrt{n} \log n)})$.
- Exponential but significantly better than the naive approach.

- **Heuristic Algorithms**

- Various heuristic methods can speed up the process.
- Exploit graph properties and use efficient data structures.

- **Average and Worst-Case Complexity**

- Average-case complexity varies with graph distribution.
- Worst-case complexity is the most critical measure.

- **Implementation Factors**

- Size of input graphs.
- Implementation details and efficiency of data structures.



Weisfeiler-Leman Test

B Weisfeiler and A Leman, "The reduction of a graph to canonical form and the algebra which appears therein," 1968



Weisfeiler-Leman Test

- Weisfeiler-Leman (WL) test is an algorithm for testing graph isomorphism.



B Weisfeiler and A Leman, "The reduction of a graph to canonical form and the algebra which appears therein," 1968

Weisfeiler-Leman Test

- Weisfeiler-Leman (WL) test is an algorithm for testing graph isomorphism.
- It operates by refining vertex labels iteratively.
- The key idea is to maintain a multiset of vertex labels.



B Weisfeiler and A Leman, "The reduction of a graph to canonical form and the algebra which appears therein," 1968

Weisfeiler-Leman Test

- Weisfeiler-Leman (WL) test is an algorithm for testing graph isomorphism.
- It operates by refining vertex labels iteratively.
- The key idea is to maintain a multiset of vertex labels.
- Efficient for graphs with bounded vertex degrees, i.e., $\deg(v_i) \leq D \quad \forall v_i \in V$, where $D \in \mathbb{I}^+$ is the upper bound.



Weisfeiler-Leman Test

- Weisfeiler-Leman (WL) test is an algorithm for testing graph isomorphism.
- It operates by refining vertex labels iteratively.
- The key idea is to maintain a multiset of vertex labels.
- Efficient for graphs with bounded vertex degrees, i.e., $\deg(v_i) \leq D \quad \forall v_i \in V$, where $D \in \mathbb{I}^+$ is the upper bound.
- Complexity: The time complexity of the WL test is $O(|V| \log |V|)$ in the worst case.
- This complexity arises from the need to sort the labels in each iteration.



Weisfeiler-Leman Test

- Weisfeiler-Leman (WL) test is an algorithm for testing graph isomorphism.
- It operates by refining vertex labels iteratively.
- The key idea is to maintain a multiset of vertex labels.
- Efficient for graphs with bounded vertex degrees, i.e., $\deg(v_i) \leq D \quad \forall v_i \in V$, where $D \in \mathbb{I}^+$ is the upper bound.
- Complexity: The time complexity of the WL test is $O(|V| \log |V|)$ in the worst case.
- This complexity arises from the need to sort the labels in each iteration.

Algorithm 1: WL-1 algorithm (Weisfeiler & Lehmann, 1968)

Input: Initial node coloring $(h_1^{(0)}, h_2^{(0)}, \dots, h_N^{(0)})$
Output: Final node coloring $(h_1^{(T)}, h_2^{(T)}, \dots, h_N^{(T)})$
 $t \leftarrow 0;$
repeat
 for $v_i \in \mathcal{V}$ **do**
 $h_i^{(t+1)} \leftarrow \text{hash} \left(\sum_{j \in \mathcal{N}_i} h_j^{(t)} \right);$
 $t \leftarrow t + 1;$
until stable node coloring is reached;

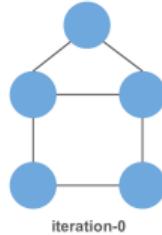
Original version formulated it as a iterative node coloring problem.



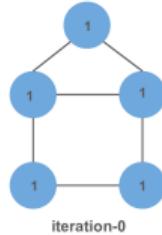
Graph Isomorphism: Example with 5 Vertices



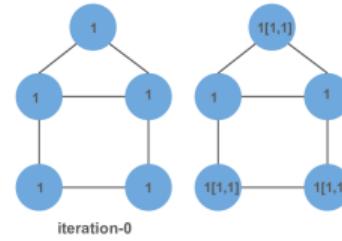
Graph Isomorphism: Example with 5 Vertices



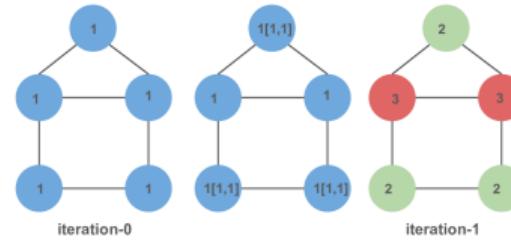
Graph Isomorphism: Example with 5 Vertices



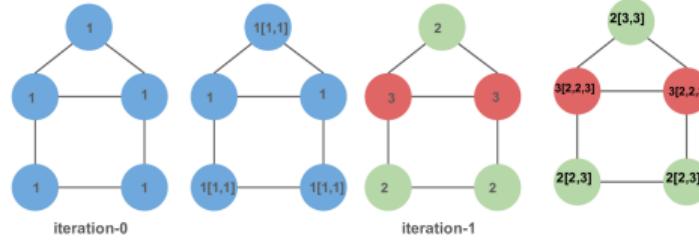
Graph Isomorphism: Example with 5 Vertices



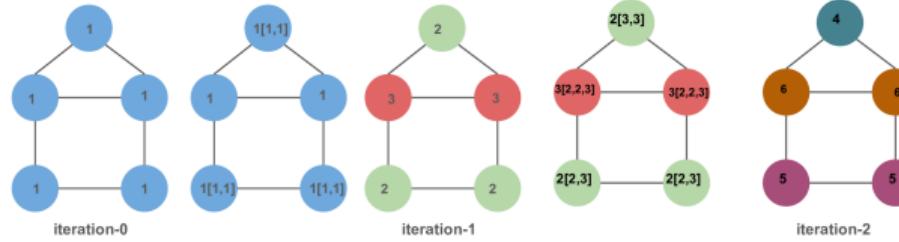
Graph Isomorphism: Example with 5 Vertices



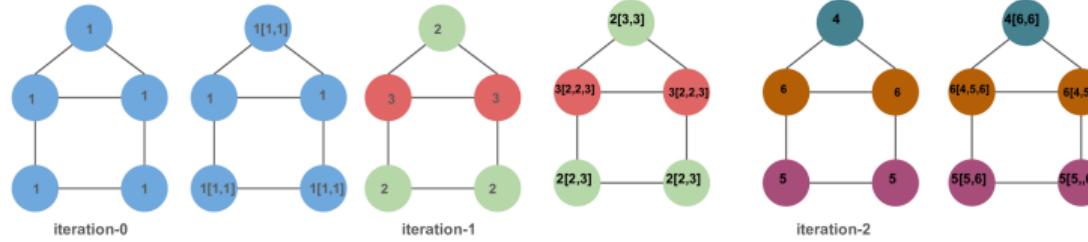
Graph Isomorphism: Example with 5 Vertices



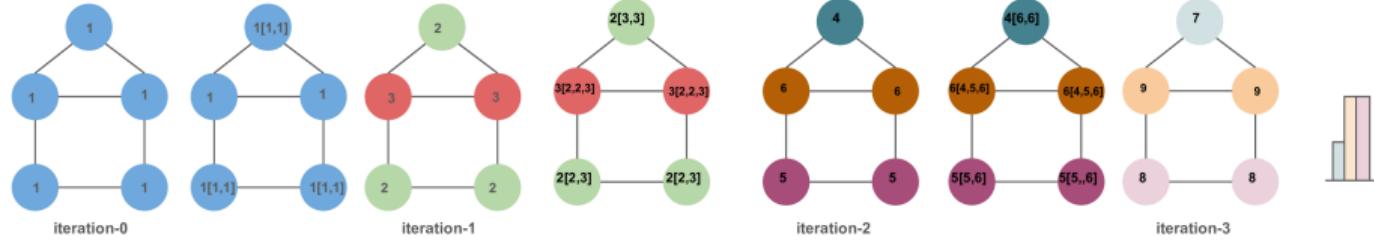
Graph Isomorphism: Example with 5 Vertices



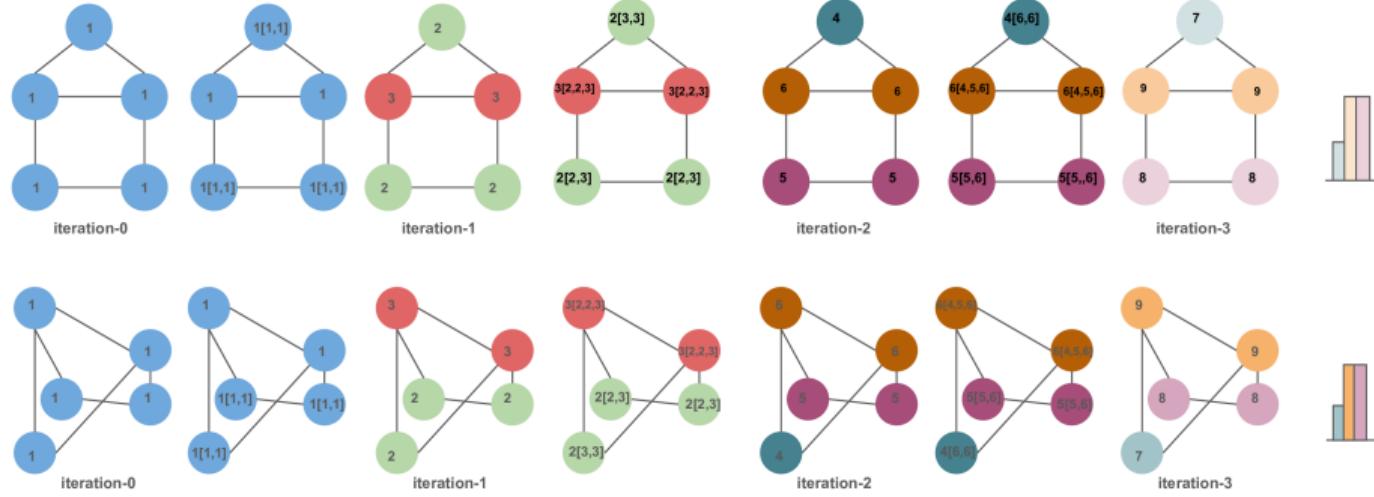
Graph Isomorphism: Example with 5 Vertices



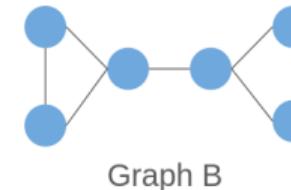
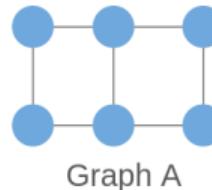
Graph Isomorphism: Example with 5 Vertices



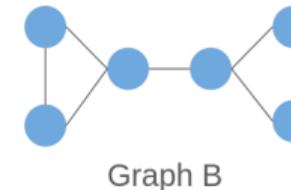
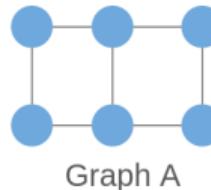
Graph Isomorphism: Example with 5 Vertices



Let's try out the 1-WL test!



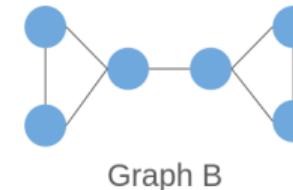
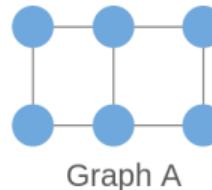
Let's try out the 1-WL test!



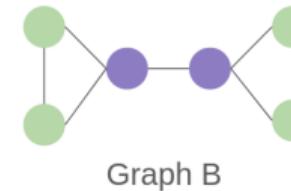
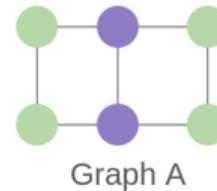
Work out the canonical graphs for the two graphs



Let's try out the 1-WL test!



Work out the canonical graphs for the two graphs

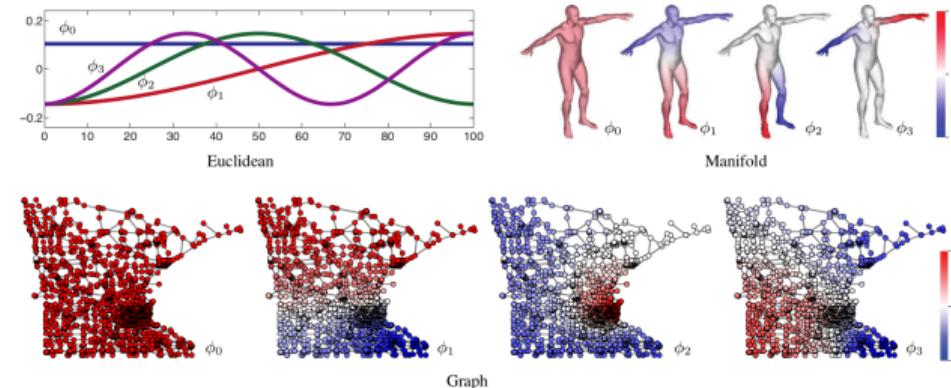


Limitations

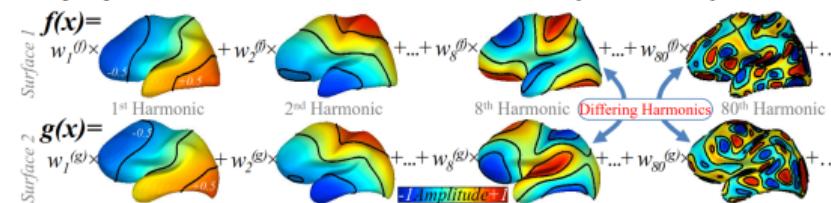
- **Non-Guarantee:** The WL test does not guarantee correctness for all instances.
- **Rare Failures:** Instances where non-isomorphic graphs have identical label multisets are rare in practice.
- **Dependency on Labeling:** The performance of the WL test depends on the initial vertex labeling.
- **Limited Application:** While versatile, the WL test may not be suitable for certain specialized cases or highly irregular graphs.



Fourier decomposition of signals



Shape Spectrum – Surface Functions can be written as composition of shape harmonics



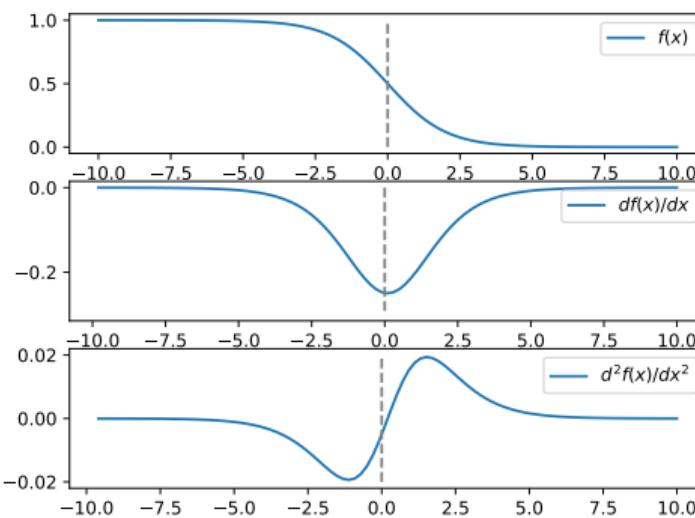
Bronstein et al. 2016. Geometric deep learning: Going beyond Euclidean data.
 Lombaert 2015. Brain Transfer: Spectral analysis of cortical surfaces and functional maps



Laplacian on different domains means (almost) the same

Laplacian: Second order differential operator

(Tells you how smoothly the function is changing in its domain)



When specified for graphs (discrete grids) → Graph Laplacian

Graph Laplacian

For a graph, \mathcal{G} with adjacency matrix $\mathbf{A} \in \mathbb{R}^{N \times N}$: the graph Laplacian is simply given as:

$$\mathbf{L} = \mathbf{D} - \mathbf{A} \quad (1)$$

$\mathbf{D} \in \mathbb{R}^{N \times N}$ is the degree matrix with $D_{ii} = \sum_j A_{ij}$

where, $\Phi = (\phi_0, \phi_1 \dots \phi_{N-1}) \in \mathbb{R}^{N \times N}$ are the orthogonal eigenfunctions forming the Fourier basis, with corresponding eigenvalues in the diagonal matrix $\Lambda = \text{diag}(\lambda_0, \lambda_1, \dots, \lambda_{N-1})$



Graph Neural Networks (GNNs): Overview

1 Lecture-1

- Motivation

 - Why Graph ML?

 - Deep Learning Basics

- Graph Learning in Spectral domain

 - ChebyNet

- Graph Learning in Spatial domain

 - Node GNNs

- Applications: Node/Graph Classification

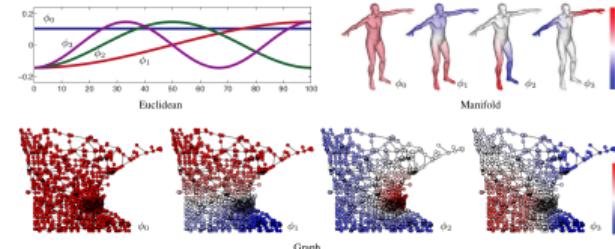


How to Perform Convolution on Graphs?

Consider a signal operating on the nodes, $x \in \mathbb{R}^N$ and a filter g_Θ parameterized by $\Theta \in \mathbb{R}^N$, the graph convolution is given as:

$$g_\Theta \star x = \quad (2)$$

where, $\Phi = (\phi_0, \phi_1 \dots \phi_{N-1}) \in \mathbb{R}^{N \times N}$ are the orthogonal eigenfunctions forming the Fourier basis



Shape Spectrum – Surface Functions can be written as composition of shape harmonics

$$\begin{aligned} f(x) &= w_1 \theta_1 x + w_2 \theta_2 x + \dots + w_8 \theta_8 x + \dots \\ \text{Surface } f &= 1^{\text{st}} \text{ Harmonic} \quad 2^{\text{nd}} \text{ Harmonic} \quad 8^{\text{th}} \text{ Harmonic} \quad 80^{\text{th}} \text{ Harmonic} \\ g(x) &= w_1 \theta_1 x + w_2 \theta_2 x + \dots + w_8 \theta_8 x + \dots \\ \text{Surface } g &= 1^{\text{st}} \text{ Harmonic} \quad 2^{\text{nd}} \text{ Harmonic} \quad 8^{\text{th}} \text{ Harmonic} \quad 80^{\text{th}} \text{ Harmonic} \end{aligned}$$

+ ... + **Differing Harmonics** + ...



Convolution on graphs in Fourier domain

Consider a signal operating on the nodes, $\mathbf{x} \in \mathbb{R}^N$ and a filter g_Θ parameterized by $\Theta \in \mathbb{R}^N$, the graph convolution is given as:

$$g_\Theta \star \mathbf{x} = (\Phi g_\Theta)(\Phi^T \mathbf{x}) \quad (3)$$

Recollect that: $\mathbf{L} = \Phi \Lambda \Phi^T$, hence $g_\Theta = \Lambda$ and $\Phi^T \mathbf{x}$ is the Fourier transform of \mathbf{x}

Turns out that $g_\Theta = \Lambda$ is not very useful

- Non-localized, non-parametric (all free parameters)
- Computationally expensive

So, use a polynomial parameterization for localized filters with:

$$g_\Theta(\Lambda) = \sum_{k=0}^{K-1} \theta_k \Lambda^k \quad (4)$$



Approximating Convolution on Graphs with Chebyshev polynomials

Recursive formulation for fast filtering:

$$g_{\Theta}(\Lambda) = \sum_{k=0}^{K-1} \theta_k \Lambda^k \approx \sum_{k=0}^{K-1} \theta_k T_k(\tilde{\Lambda}) \quad (5)$$

with $\tilde{\Lambda} = \frac{2}{\lambda_{\max}} \Lambda - \mathbf{I}_N \in [-1, 1]$ and Chebyshev coefficients $\Theta \in \mathbb{R}^K$.

ChebyNet

Incorporating the recursive approximation yields the popular spectral graph convolution method:
ChebyNet

$$g_{\Theta} \star \mathbf{x} = (\Phi g_{\Theta})(\Phi^T \mathbf{x}) \approx \sum_{k=0}^{K-1} \theta_k T_k(\tilde{\mathbf{L}}) \mathbf{x} \quad (6)$$

with $\tilde{\mathbf{L}} = \frac{2}{\lambda_{\max}} \mathbf{L} - \mathbf{I}_N$.

Chebyshev polynomial $T_k(y)$ of order k is given by the recurrence: $T_k(y) = 2yT_{k-1}(y) - T_{k-2}(y)$ with $T_0 = 1$, $T_1 = y$
 Defferrard et al. (2016). Convolutional neural networks on graphs with fast localized spectral filtering.



Overview of ChebyNet: CNN for Graphs (but in spectral domain)

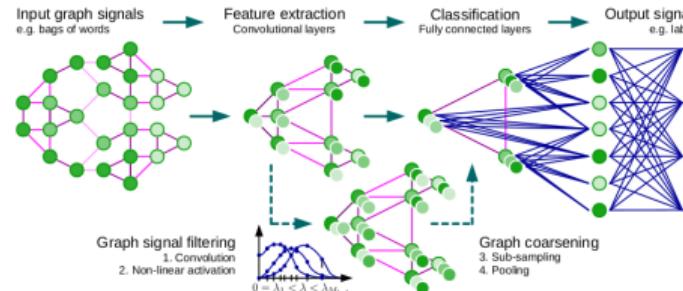
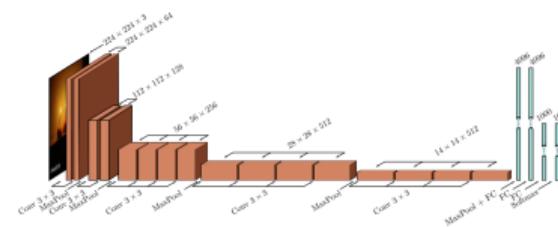


Figure 1: Architecture of a CNN on graphs and the four ingredients of a (graph) convolutional layer.

Spectral CNNs mimic standard CNNs

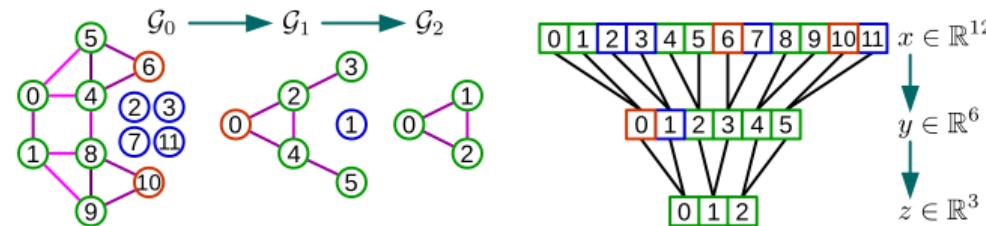


Defferrard et al. Convolutional Neural Networks on Graphs with Fast Localized Spectral Filtering. 2016



How to Perform Pooling on Graphs?

Graph Clustering is NP-hard



Based on a greedy algorithm (Graculus), accelerated by using balanced binary trees using fake nodes.

Defferrard et al. Convolutional Neural Networks on Graphs with Fast Localized Spectral Filtering. 2016
 I. Dhillon, Y. Guan, and B. Kulis. Weighted Graph Cuts Without Eigenvectors: A Multilevel Approach. 2007.
 Liu et al. Graph Pooling for Graph Neural Networks: Progress, Challenges, and Opportunities. 2023



Graph Neural Networks (GNNs): Overview

1 Lecture-1

- Motivation
 - Why Graph ML?
 - Deep Learning Basics
- Graph Learning in Spectral domain
 - ChebyNet
- **Graph Learning in Spatial domain**
 - Node GNNs
- Applications: Node/Graph Classification



Further approximation of ChebyNet yields the class of Node GNNs

Recall that, from ChebyNet, $g_\Theta \star \mathbf{x} = (\Phi g_\Theta)(\Phi^T \mathbf{x}) \approx \left(\Phi \sum_{k=0}^{K-1} \theta_k T_k(\tilde{\mathbf{L}}) \right) (\Phi^T \mathbf{x}) = \sum_{k=0}^{K-1} \theta_k T_k(\tilde{\mathbf{L}}) \mathbf{x}$,
 With $K = 2$, $\lambda_{max} = 2$, the graph convolution operation becomes:

$$g_\Theta \star \mathbf{x} \approx \theta_0 \mathbf{x} + \theta_1 \tilde{\mathbf{L}} \mathbf{x} \quad (7)$$

$$= \theta_0 \mathbf{x} + \theta_1 (\mathbf{L} - \mathbf{I}_N) \mathbf{x} \quad (8)$$

$$= \theta (\mathbf{I}_N - \mathbf{D}^{-1/2} \mathbf{A} \mathbf{D}^{-1/2}) \mathbf{x} \quad (9)$$

$$g_\Theta \star \mathbf{x} \approx \theta (\tilde{\mathbf{D}}^{-1/2} \tilde{\mathbf{A}} \tilde{\mathbf{D}}^{-1/2} \mathbf{x}) \quad (10)$$

with $\theta_0 = -\theta_1$, and $\tilde{\mathbf{A}} = \mathbf{A} + \mathbf{I}$. More generally, for input $\mathbf{X} \in \mathbb{R}^{N \times F}$ and weight matrix $\Theta \in \mathbb{R}^{F \times L}$:

$$\mathbf{H} = \sigma(\hat{\mathbf{A}} \mathbf{X} \Theta) \quad (11)$$

Stacking multiple of these layers with non-linearities yields the class of node GNNs³!

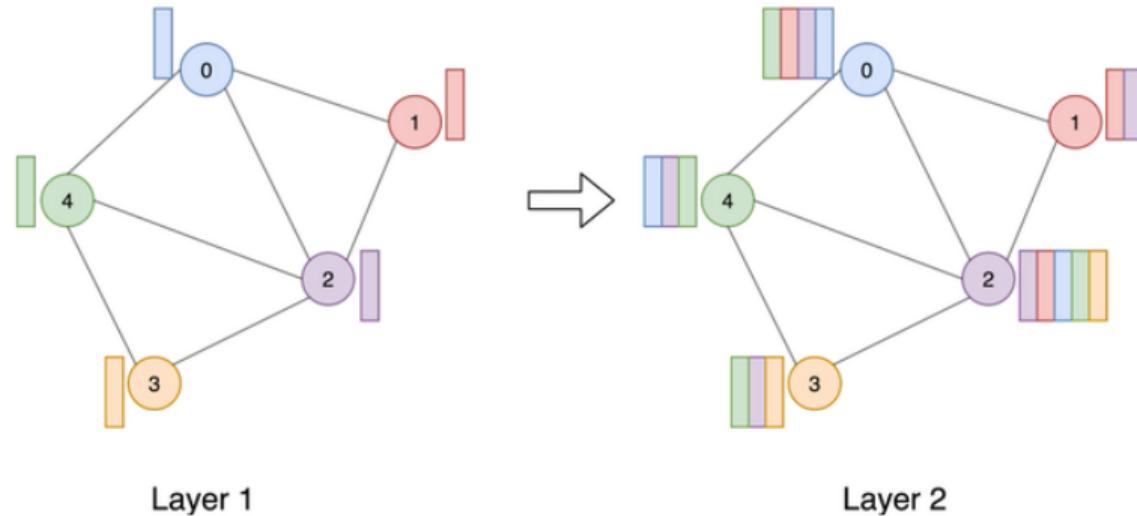
³Kipf, Thomas N., and Max Welling. "Semi-supervised classification with graph convolutional networks." (2016)

Note: Chebyshev polynomial $T_k(y)$ of order k is given by the recurrence: $T_k(y) = 2yT_{k-1}(y) - T_{k-2}(y)$ with $T_0 = 1$, $T_1 = y$

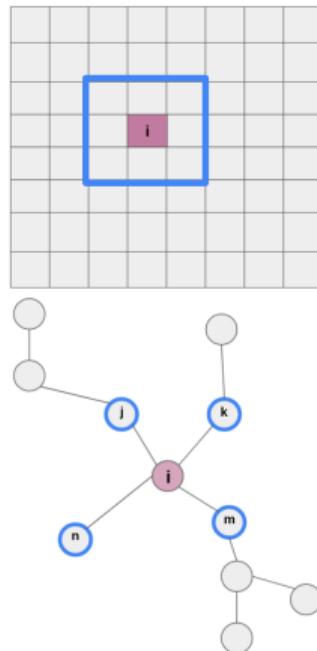


GNNs update states of nodes based on neighbourhood

Graphically, $\mathbf{H} = \sigma(\hat{\mathbf{A}}\mathbf{X}\Theta)$ means:



GNNs when seen from a node's point of view



- For a node $i \in \mathcal{V}$ with neighbours \mathcal{N}_i the GNN operation in layer- ℓ is given as:

$$\mathbf{h}_i^{(\ell)} = \sigma \left(\frac{\Theta^{(\ell)}}{|\mathcal{N}_i|} \sum_{j \in \mathcal{N}_i} \mathbf{h}_j^{(\ell-1)} \right) \quad (12)$$

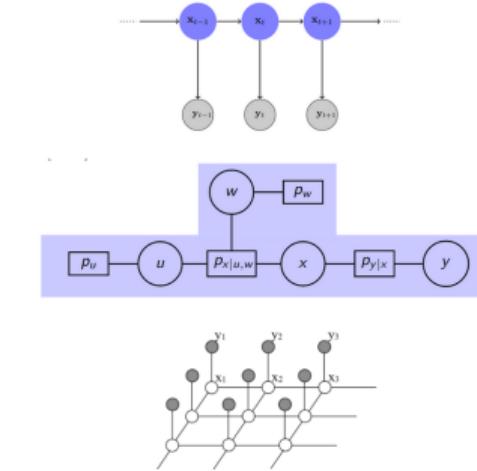
with $\mathbf{h}_i^{(0)} = \mathbf{x}_i$ and $\Theta^{(\ell)}$ are trainable parameters.

- Aggregation of transformed neighbouring node features
- M-layered GNN provides information from M-hops away!
- A form of learnable message passing!



Rich classes of message passing algorithms exist for inference on PGMs

- Kalman filter for Markov chains
- Sum-product algorithm for factor graphs
- Loopy belief propagation for graphs with cycles
- Mean field approximation for Markov random fields

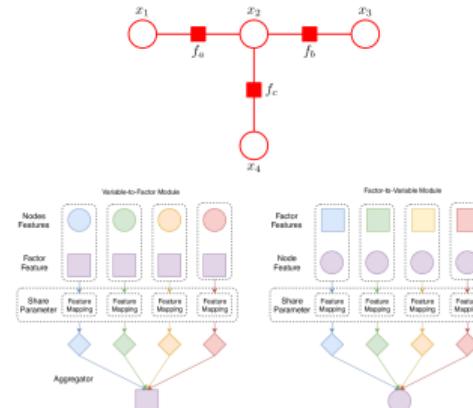


Can GNNs be interpreted as generalizations of these algorithms?



Message Passing Neural Networks

- Spectral and Spatial to notion of generalized message passing
- Close connections to probabilistic graphical models (PGMs)
- Similarities with classic methods in PGMs like max-product algorithm/ belief propagation



Gilmer et al. Neural Message Passing for Quantum Chemistry. 2017.
 Bishop. Pattern Recognition and Machine Learning. 2006
 Zhang et al. Factor Graph Neural Network. 2020.



Graph Neural Networks (GNNs): Overview

1 Lecture-1

- Motivation
 - Why Graph ML?
 - Deep Learning Basics
- Graph Learning in Spectral domain
 - ChebyNet
- Graph Learning in Spatial domain
 - Node GNNs
- Applications: Node/Graph Classification



Node Classification using GCNs

Observed data $\mathbf{G} : \mathbf{X} \in \mathbb{R}^{N \times F}, \mathbf{A} \in \mathbb{R}^{N \times N}$

Node-level Labels $\mathbf{Y} : \{0, 1\}^N$

Node Classifier $g_\theta(\cdot) : \mathbf{G} \rightarrow \mathbf{Y}$

L-layered GCN for node classification

$$\mathbf{H}^{(\ell)} = \sigma(\hat{\mathbf{A}} \mathbf{H}^{(\ell-1)} \boldsymbol{\Theta}^{(\ell)}) \quad (13)$$

where, $\mathbf{H}^{(0)} = \mathbf{X}$ and $\mathbf{H}^{(L)} = \mathbf{Y}$

Network weights $\boldsymbol{\Theta}$ are trained in a supervised manner by optimizing an appropriate loss function:
 $\mathcal{L}(\mathbf{Y}, \hat{\mathbf{Y}})$.



Graph Classification using GCNs

Observed data: $\mathbf{G} = \{\mathbf{G}_1, \mathbf{G}_2, \dots, \mathbf{G}_N\} : \mathbf{G}_i : \mathbf{X}_i, \mathbf{A}_i, \mathbf{E}_i$

Labels $\mathbf{Y} = \{\mathbf{y}_1, \mathbf{y}_2, \dots, \mathbf{y}_N\} : \mathbf{y}_i \in [0, C - 1]^N$

Graph Classifier $g_\theta(\cdot) : \mathbf{G} \rightarrow \mathbf{Y}$

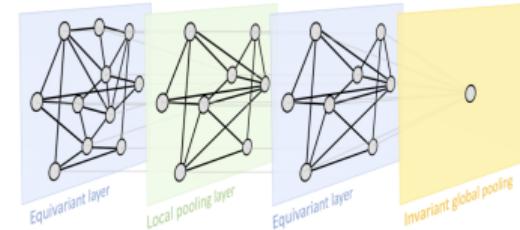


Figure 8: Geometric Deep Learning blueprint, exemplified on a graph. A typical Graph Neural Network architecture may contain permutation equivariant layers (computing node-wise features), local pooling (graph coarsening), and a permutation-invariant global pooling layer (readout layer).

L-layered GCN for node classification

$$\mathbf{H}^{(\ell)} = \sigma(\hat{\mathbf{A}} \mathbf{H}^{(\ell-1)} \Theta^{(\ell)}) \quad (14)$$

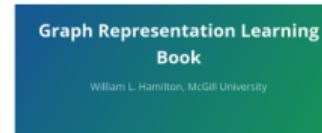
$$\hat{\mathbf{Y}} = \text{Aggregate}(\mathbf{H}^L) \quad (15)$$

where, $\mathbf{H}^{(0)} = \mathbf{X}$

Trained in a fully supervised manner by optimizing an appropriate loss function: $\mathcal{L}(\mathbf{Y}, \hat{\mathbf{Y}})$.

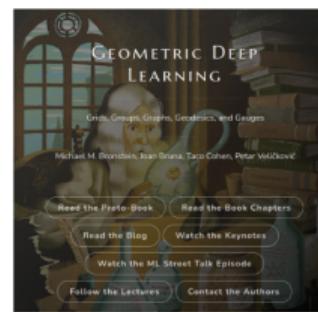


Resources



The field of graph representation learning has grown at an incredible (and sometimes unwieldy) pace over the past seven years, transforming from a small subset of researchers working on a relatively niche topic to one of the fastest growing sub-areas of deep learning.

This book is my attempt to provide a brief but comprehensive introduction to graph representation learning, including methods for embedding graph data, graph neural networks, and deep generative models of graphs.



https://www.cs.mcgill.ca/~wlh/grl_book/
<https://geometricdeeplearning.com/>
<https://ogb.stanford.edu/>
<https://pytorch-geometric.readthedocs.io/en/latest/>



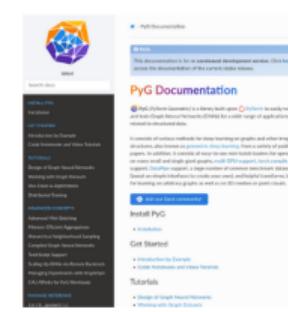
The Open Graph Benchmark (OGB) is a collection of realistic, large-scale, and diverse benchmarks datasets for machine learning on graphs. OGB provides a collection of benchmarks, including citation networks, and split using the [ML4D Data License](#). The model performance can be evaluated using the [OGBL](#) evaluation or a performance monitor.



Realistic datasets
Independently curated set of challenging and diverse benchmarks for graph learning. OGB provides a collection of benchmarks, including citation networks, and split using the [ML4D Data License](#).

Flexible data loaders
ML4D provides flexible data loaders and data processing pipelines for loading and processing graph-structured data such as [PyTorch Geometric](#) and [GNNExplainer](#).

Unified evaluation
ML4D provides standard evaluation metrics and evaluation tools for comparing different methods. OGB uses [PyTorch Geometric](#) to implement the state-of-the-art.



Summary

- Convolutions on graphs can be approximated in spectral domain
- ChebyNet uses a polynomial of spectral filter approximated with Chebyshev polynomials
- First order approximation to ChebyNet yields spatial graph convolutions
- Message passing networks can be used learn *useful* embeddings for downstream tasks
- Simple GCNs for node classification; can be extended to graph classification models
- Community started out with convolution on graphs (GCNs)
- Currently, primarily viewed as learnable message passing (GNNs)

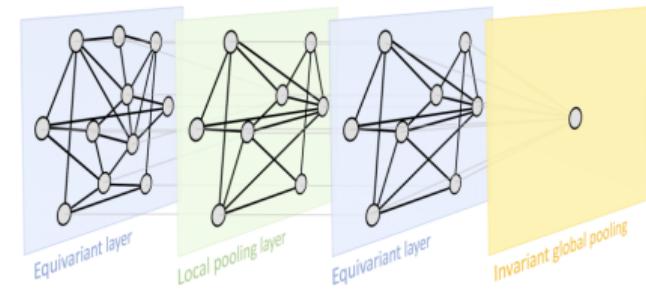


Figure 8: Geometric Deep Learning blueprint, exemplified on a graph. A typical Graph Neural Network architecture may contain permutation equivariant layers (computing node-wise features), local pooling (graph coarsening), and a permutation-invariant global pooling layer (readout layer).

